

Technical Document #120: Machine Learning - Comprehensive Overview

Technical Document #120: Machine Learning - Comprehensive Overview

Gradient descent is an optimization algorithm used to minimize the cost function in machine learning models. It iteratively adjusts parameters in the direction of steepest descent.

Dimensionality reduction techniques like PCA and t-SNE help visualize high-dimensional data and can improve model performance by removing irrelevant features.

Support vector machines (SVMs) are supervised learning models that analyze data for classification and regression analysis. They find the hyperplane that best separates different classes.

The bias-variance tradeoff is a fundamental concept in machine learning. High bias leads to underfitting while high variance leads to overfitting.

Overfitting occurs when a model learns the training data too well, including noise and outliers, leading to poor generalization on unseen data. Regularization techniques help prevent this.

Artificial intelligence continues to reshape industries from healthcare to finance. Machine learning models can now perform tasks that were once thought to require human intelligence.

Federated learning trains models across decentralized data sources without sharing raw data. It enables privacy-preserving machine learning.

Cybersecurity threats continue to evolve in complexity and scale. Organizations must invest in robust security measures to protect sensitive data.

Key Takeaways:

- Overfitting occurs when a model learns the training data too well, including noise and outliers, leading to poor generalization on unseen data.
- The bias-variance tradeoff is a fundamental concept in machine learning.
- Feature selection is the process of selecting a subset of relevant features for use in model construction.