Technical Document #636: Retrieval Augmented Generation - Comprehensive Overvi

Technical Document #636: Retrieval Augmented Generation - Comprehensive Overview

Hallucination detection identifies when LLMs generate factually incorrect information. It is critical for maintaining trust in RAG systems.

Re-ranking models score retrieved documents for relevance to the query. They improve the quality of context provided to the generator.

Retrieval-Augmented Generation (RAG) combines information retrieval with text generation. It grounds LLM responses in factual, retrieved documents.

Data-driven decision making has become essential for modern businesses. Organizations that leverage data effectively gain a significant competitive advantage.

Remote work technologies have transformed the workplace. Collaboration tools and cloud services enable distributed teams to work effectively from anywhere.

Artificial intelligence continues to reshape industries from healthcare to finance. Machine learning models can now perform tasks that were once thought to require human intelligence.

Key Takeaways:

- Context window optimization selects the most relevant information to fit within the LLM's context limit.
- Vector databases store dense embeddings for semantic search.
- Retrieval-Augmented Generation (RAG) combines information retrieval with text generation.