Technical Document #687: Data Engineering - Comprehensive Overview

Technical Document #687: Data Engineering - Comprehensive Overview

Schema evolution handles changes in data structure over time. Schema registries manage schema versions and ensure backward compatibility.

ETL (Extract, Transform, Load) is a process that extracts data from source systems, transforms it to fit business rules, and loads it into a target data store.

Data warehouses store structured data optimized for analytical queries. They use columnar storage and support complex aggregations.

Partitioning divides data into smaller, more manageable pieces. It improves query performance by reducing the amount of data scanned.

Data quality ensures data is accurate, complete, and consistent. Data validation, cleansing, and monitoring are essential components.

Data lineage tracks the origin and movement of data through systems. It is essential for debugging, compliance, and impact analysis.

Computer vision applications are becoming ubiquitous, from facial recognition to autonomous driving. Deep learning has enabled breakthroughs in visual understanding.

Autonomous vehicles use a combination of sensors, AI, and mapping to navigate without human intervention. Self-driving technology is rapidly advancing.

Federated learning trains models across decentralized data sources without sharing raw data. It enables privacy-preserving machine learning.

Voice assistants like Alexa and Siri use speech recognition and natural language understanding. They have become integral parts of smart home ecosystems.

Cybersecurity threats continue to evolve in complexity and scale. Organizations must invest in robust security measures to protect sensitive data.

Key Takeaways:

- Schema evolution handles changes in data structure over time.
- Data pipelines automate the movement and transformation of data between systems.
- Data warehouses store structured data optimized for analytical queries.