

# Technical Document #754: Retrieval Augmented Generation - Comprehensive Overview

## Technical Document #754: Retrieval Augmented Generation - Comprehensive Overview

Retrieval-Augmented Generation (RAG) combines information retrieval with text generation. It grounds LLM responses in factual, retrieved documents.

Semantic search finds documents based on meaning rather than exact keywords. It uses embedding similarity to match queries with relevant content.

Hallucination detection identifies when LLMs generate factually incorrect information. It is critical for maintaining trust in RAG systems.

Evaluation metrics for RAG include faithfulness, answer relevance, and context recall. They measure the quality of the retrieval-generation pipeline.

Vector databases store dense embeddings for semantic search. Pinecone, Weaviate, and ChromaDB are popular vector database solutions.

Artificial intelligence continues to reshape industries from healthcare to finance. Machine learning models can now perform tasks that were once thought to require human intelligence.

Digital twins are virtual replicas of physical systems. They enable simulation and optimization without risking real-world resources.

Remote work technologies have transformed the workplace. Collaboration tools and cloud services enable distributed teams to work effectively from anywhere.

Voice assistants like Alexa and Siri use speech recognition and natural language understanding. They have become integral parts of smart home ecosystems.

Quantum computing promises to solve problems that are intractable for classical computers. It has potential applications in cryptography, drug discovery, and optimization.

Supply chain management leverages technology to optimize logistics and distribution. Real-time tracking and predictive analytics improve efficiency.

### Key Takeaways:

- Hallucination detection identifies when LLMs generate factually incorrect information.
- Hybrid search combines keyword-based and semantic search.
- Vector databases store dense embeddings for semantic search.