

Technical Document #787: Computer Vision - Comprehensive Overview

Technical Document #787: Computer Vision - Comprehensive Overview

Image classification assigns a label to an image from a fixed set of categories. CNNs like ResNet and EfficientNet achieve superhuman performance on benchmark datasets.

Face recognition identifies or verifies faces in images. DeepFace and FaceNet achieve near-human accuracy on face verification tasks.

Visual question answering (VQA) answers questions about images. It requires both visual understanding and language comprehension.

Semantic segmentation classifies each pixel in an image into a category. It is used in autonomous driving and medical image analysis.

Image super-resolution reconstructs high-resolution images from low-resolution inputs. GAN-based methods produce photorealistic upscaled images.

Object detection identifies and localizes objects in images. YOLO and Faster R-CNN are popular real-time object detection architectures.

Voice assistants like Alexa and Siri use speech recognition and natural language understanding. They have become integral parts of smart home ecosystems.

Augmented reality overlays digital information on the physical world. It has applications in training, maintenance, and entertainment.

Computer vision applications are becoming ubiquitous, from facial recognition to autonomous driving. Deep learning has enabled breakthroughs in visual understanding.

Autonomous vehicles use a combination of sensors, AI, and mapping to navigate without human intervention. Self-driving technology is rapidly advancing.

Key Takeaways:

- Image super-resolution reconstructs high-resolution images from low-resolution inputs.
- Video understanding analyzes temporal dynamics in video sequences.
- Semantic segmentation classifies each pixel in an image into a category.